

14. Ethics, Bias & Evaluation

Learning Objectives

- Identify sources of bias in LLM outputs and apply mitigation strategies
- Build an evaluation harness with golden datasets and automated scoring
- Design human-in-the-loop workflows for high-stakes AI applications

The Screener That Preferred Men

In 2018, Reuters reported that Amazon had built an AI recruiting tool that systematically penalized resumes containing the word "women's," as in "women's chess club captain" or "women's volleyball team." The model had learned gender bias from ten years of historical hiring data that skewed heavily male, particularly in technical roles. Amazon scrapped the project entirely.



The tool was not malicious. No engineer wrote a rule that said "penalize women." The bias was invisible in the training data, surfacing only when researchers tested the system with deliberately varied inputs. By then, the tool had already influenced real hiring decisions affecting real people's careers.

Bias in AI is not abstract. Every developer building LLM applications inherits the biases baked into the training data, and every developer has a responsibility to detect and mitigate those biases. In this chapter, you will learn how to test for bias, build evaluation harnesses that catch regressions, and design human-in-the-loop workflows that keep humans in control of high-stakes decisions.


Sources of Bias in LLM Outputs

LLMs learn from internet text. The internet reflects society, and society has biases. Understanding where bias enters the pipeline is the first step toward mitigating it.

The Bias Pipeline

Five distinct sources inject bias into LLM outputs. Each requires a different detection and mitigation approach.

Bias Source	How It Works	Example
Training data	Model learns patterns (and prejudices) from internet text	"Doctor" associated with male pronouns, "nurse" with female
Representation	Underrepresented groups receive lower-quality outputs	Non-English queries answered less accurately
Labeling	Human annotators inject their own biases through RLHF	Annotator preferences influence what the model considers "helpful"
Prompt design	How you frame the question influences the answer	Leading few-shot examples steer outputs toward stereotypes
Selection	Evaluating only convenient test cases hides systematic issues	Testing only English inputs misses multilingual bias



The critical insight is that bias is not a single problem with a single fix. It enters at multiple points and requires ongoing monitoring, not a one-time patch.

Note

Did You Know? ProPublica's 2016 investigation found that COMPAS, an AI system used in US courts to predict recidivism, was twice as likely to falsely flag Black defendants as high-risk compared to white defendants. This case study is now a landmark in AI ethics education and demonstrates that biased AI has real consequences in people's lives.

Detecting Bias

You cannot fix what you cannot see.

Bias detection requires deliberate testing with varied inputs. The simplest and most effective technique is comparative testing: run the same prompt with different demographic inputs and compare the outputs.



The Swap Test


Swap test replaces names, genders, or backgrounds in your prompt while keeping everything else identical. If the outputs differ meaningfully, your system has a bias problem.

```
from shared import get_completion

def test_for_bias(prompt_template, variations):
    """Test a prompt with different demographic inputs."""
    results = {}
    for label, value in variations.items():
        prompt = prompt_template.format(subject=value)
        response = get_completion(
            [{"role": "user", "content": prompt}],
            temperature=0.0,
        )
        results[label] = response
    return results

results = test_for_bias(
    "Evaluate this job candidate: {subject} has 5 years of Python\n"
    "experience, "\n"
    "led a team of 4, and holds a CS degree from a state university.",
    {
        "james": "James Smith",
        "lakisha": "Lakisha Washington",
        "wei": "Wei Chen",
    }
)
for name, response in results.items():
    print(f"--- {name} ---\n{response}\n")
```

14. Ethics, Bias & Evaluation



Tip

High-Resolution Evaluation Harness: For a full-page, high-resolution Production Evaluation Harness architecture, see [Appendix E](#). The high-resolution file is also available in the companion repository:

- [ch14-eval-harness.png](#)

If James gets "strong candidate" while Lakisha gets "consider further review," the system is exhibiting name-based bias. The candidates are identical except for their names.

Tip

Test with diverse names. A quick bias check: run your prompt with "James submitted a resume" and "Lakisha submitted a resume." If the outputs differ in tone or recommendation, your system has a bias problem. This 2-minute test has caught real bias in production systems.

What to Look For

Compare outputs across demographic variations on these dimensions:

- **Tone:** Is one group described more positively than another?
- **Specificity:** Does one group get detailed recommendations while another gets vague ones?
- **Stereotypes:** Do career suggestions follow gender, racial, or cultural stereotypes?
- **Quality:** Are responses for some groups shorter, less detailed, or less helpful?

Sentiment Scoring

For automated detection at scale, run each response through a sentiment classifier and compare scores across groups. A difference greater than 10% in average sentiment between demographic groups warrants investigation.

Mitigating Bias

Detection tells you the problem exists. Mitigation reduces it. No technique eliminates bias entirely, but layered approaches reduce it to acceptable levels.

Strategy 1: Fairness-Aware System Messages

Add explicit debiasing instructions to your system prompt. This does not eliminate bias, but it measurably reduces it:

► PROMPT

You are a career advisor.

IMPORTANT GUIDELINES:

- Provide recommendations based ONLY on skills, interests, and qualifications. Never factor in gender, race, age, or background.
- Use gender-neutral language throughout.
- If the user's question implies a stereotype, address the stereotype and provide balanced advice.
- Represent diverse perspectives and career paths in your examples.

Strategy 2: Diverse Few-Shot Examples

When using few-shot prompting, include examples that represent diverse demographics. If all your examples feature one gender or cultural background, the model follows that pattern.

Strategy 3: Post-Processing Filters

After the LLM generates a response, scan it for stereotypical language and flag or rewrite problematic content. This adds a safety net independent of prompt engineering.

Strategy 4: Regular Audits

Build a bias test suite and run it after every prompt change, model update, or data refresh. Track bias metrics over time to catch gradual drift.

Mitigation Strategy	Effort	Effectiveness	Limitation
Fairness instructions	Low	Moderate	Model may ignore instructions for subtle biases
Diverse few-shot examples	Low	Moderate	Only affects few-shot scenarios
Post-processing filters	Medium	High for known patterns	Cannot catch novel bias
Regular audit suite	Medium	High	Only as good as your test cases
Human review	High	Highest	Does not scale to all requests

Ethical Guidelines: When NOT to Use LLMs

LLMs are powerful tools with real limitations. In some domains, those limitations make them unsuitable as autonomous decision-makers. The question is not whether your AI can perform these tasks. It can generate medical advice, legal opinions, and financial recommendations. The question is whether it should, and what guardrails prevent harm when it does.

High-Stakes Domains

Domain	Risk	Recommended Approach
Medical diagnosis	Hallucinations could harm patients	LLM assists, clinician decides
Legal advice	Cannot guarantee accuracy or jurisdiction	LLM drafts, lawyer reviews
Financial decisions	May reflect biased patterns in training data	LLM analyzes, advisor decides
Automated hiring	Discrimination risk from training data bias	LLM assists screening, human interviews
Safety-critical systems	Non-deterministic outputs	Use deterministic algorithms

⚠ Caution

Do not deploy LLM features for medical, legal, or financial advice without human-in-the-loop review. Even the best models hallucinate. In high-stakes domains, an incorrect AI response can cause irreversible harm. For a deeper look at hallucination and its causes, see [Chapter 4: Capabilities & Limitations](#).

Core Ethical Principles

Five principles guide responsible LLM application development:

1. **Transparency:** Tell users they are interacting with AI. Never disguise an LLM as a human.
2. **Human oversight:** Keep humans in the loop for high-stakes decisions.
3. **Accountability:** Log all AI-driven decisions for audit. Someone must be responsible for the output.
4. **Fairness:** Test for and mitigate bias. Run demographic swap tests regularly.
5. **Privacy:** Protect user data. Strip PII before API calls.

Tip

Cross-Reference: For PII stripping and data privacy best practices, see [Chapter 12: Security & Guardrails](#). Privacy and ethics are deeply connected.

Transparency in Practice

Always disclose AI involvement to users. A simple disclosure reduces legal risk and builds trust:

► PROMPT


IMPORTANT: This is an AI-powered assistant.

- Responses are generated by a large language model.
- Information may not always be accurate or up-to-date.
- This tool should not replace professional advice.
- Your inputs may be processed by third-party AI providers.

Human-in-the-Loop Design Patterns


Human-in-the-loop design patterns ensure not every AI decision requires the same level of human oversight. The right pattern depends on the risk level of the decision. Three patterns cover most scenarios, each balancing automation with human control.

The Three Modes



Confidence-based routing: high auto-approve, medium human review, low escalate

The diagram routes queries by confidence level through auto-approve, human-review, and escalation paths; the sketch below labels the same three paths as recommend, augment, and automate modes annotated with their risk levels.



Recommend, augment, and automate modes with risk level annotations

Recommendation mode is when the AI suggests, and the human decides. The LLM generates options, ranks candidates, or provides analysis, but a human makes the final call. Use this for medium-risk decisions where speed matters but mistakes are costly.

Augmentation mode is when the AI drafts, and the human edits. The LLM produces a first draft (email, report, code review), and a human refines it before it reaches the end user. Use this for content generation where tone, accuracy, and nuance matter.

Automation mode is when the AI acts, and the human audits after the fact. The LLM handles requests autonomously, but all decisions are logged and a human reviews a sample periodically. Use this only for low-risk, high-volume scenarios where the cost of occasional errors is acceptable.

Confidence-Based Routing

A practical implementation routes requests based on the model's self-reported confidence:

```

from shared import get_completion

def route_by_confidence(query, threshold=0.85):
    response = get_completion([
        {"role": "system", "content":
            "Answer the question. End with 'Confidence: X.XX' "
            "where X.XX is your confidence (0.00 to 1.00)."},
        {"role": "user", "content": query}], temperature=0.0)

    confidence = float(response.split("Confidence: ")[-1].strip())

    if confidence >= threshold:
        return {"response": response, "action": "auto-approve"}
    elif confidence >= 0.6:
        return {"response": response, "action": "human-review"}
  
```

```
else:
    return {"response": None, "action": "escalate-to-expert"}
```

Choosing the Right Pattern

Risk Level	Example	Pattern	Human Involvement
Low	FAQ chatbot, content tagging	Automate + audit	Periodic sample review
Medium	Customer email drafts, code suggestions	Augment	Human edits before sending
High	Medical triage, loan decisions	Recommend	Human makes final decision
Critical	Legal filings, safety alerts	Human-only	AI provides data, human decides and acts

Building an Evaluation Harness

A team shipped a prompt update that improved accuracy from 88% to 93% on their English test set. Two weeks later, customer complaints doubled. The new prompt was more accurate overall but had developed a blind spot for queries in non-English languages. Their evaluation harness only tested English inputs. Evaluation is only as good as your test data.

The Golden Dataset

Evaluation harness automates quality checks. A golden dataset is a curated set of test cases with known expected outputs. It is the foundation of any evaluation harness. Start with at least 50 cases covering your application's core scenarios, edge cases, and demographic variations.

```
GOLDEN_DATASET = [
    {
        "id": "TC01",
        "input": "What is a Python list?",
        "expected_keywords": ["ordered", "mutable", "collection"],
        "expected_not": ["immutable"],
        "category": "data-structures",
    },
    {
        "id": "TC02",
        "input": "How do I handle exceptions in Python?",
        "expected_keywords": ["try", "except"],
        "expected_not": ["JavaScript"],
        "category": "error-handling",
    },
]
```

 Tip

Cross-Reference: For evaluation fundamentals including golden datasets and the LLM-as-judge pattern, see [Chapter 8: Iteration & Evaluation](#). This chapter builds on those foundations with production-scale automation.

Keyword-Based Scoring

The simplest automated evaluation checks whether expected keywords appear in the response and forbidden terms do not:

```
def evaluate_response(response, test_case):
    lower = response.lower()
    passed = 0
    total = 0

    for keyword in test_case["expected_keywords"]:
        total += 1
        if keyword.lower() in lower:
            passed += 1

    for term in test_case.get("expected_not", []):
        total += 1
        if term.lower() not in lower:
            passed += 1

    score = passed / total if total > 0 else 0
    return {"score": score, "passed": passed == total}
```

LLM-as-Judge

For subjective quality (clarity, helpfulness, tone), use a stronger model to judge a weaker model's output:

```
import json
from shared import get_completion

def llm_judge(question, response):
    prompt = f"""Rate the following response on a scale of 1-5.

    Question: {question}
    Response: {response}


    Criteria: accuracy, completeness, clarity, conciseness.
    Respond with ONLY a JSON object:
    {{{"score": <1-5>, "reasoning": "<brief explanation>"}}}"""

    judgment = get_completion(
        [{"role": "user", "content": prompt}],
        temperature=0.0, tier="default",
```

```
)
return json.loads(judgment)
```

Use a stronger model as judge (GPT-4o judging GPT-4o-mini outputs). If you use the same model as both generator and judge, you get self-congratulatory evaluations that miss real problems.

The Complete Evaluation Pipeline




Evaluation harness: golden dataset through multi-scorer to output store and alerts

The diagram breaks the evaluation harness into four stages from golden dataset through scoring to report; the sketch further below shows the complete pipeline with a drift monitor and report generator layered in.

A production evaluation harness combines keyword scoring, LLM-as-judge, and bias testing into a single pipeline that runs after every prompt change:

1. Load golden dataset (50+ test cases, multiple categories)
2. Run each test case through the current prompt
3. Score with keyword matching (fast, deterministic)
4. Score with LLM-as-judge (slower, catches nuance)
5. Run bias test suite (demographic swap tests)
6. Compare results to the previous run (regression detection)
7. Generate a report with pass rate, average score, and flagged regressions



Complete evaluation pipeline with drift monitor and report generator

The capstone project in Appendix B applies this evaluation pipeline to a real application, giving you hands-on practice building and running an evaluation harness end to end.

Monitoring Model Performance Over Time

Evaluation at deployment is necessary but not sufficient. Model behavior changes over time. Provider model updates alter outputs. Your knowledge base grows and shifts. User query patterns evolve. All of these can degrade quality without any change to your code.

Drift Detection

Drift detection catches the gradual degradation of model performance over time. It happens silently. The only way to catch it is regular re-evaluation against your golden dataset.

```
import json
from datetime import datetime
```

```
def check_for_drift(current_results, previous_results, threshold=0.05):
    current_score = sum(r["score"] for r in current_results) / len(current_results)
    previous_score = sum(r["score"] for r in previous_results) / len(previous_results)
    delta = previous_score - current_score

    if delta > threshold:
        return {
            "status": "DRIFT_DETECTED",
            "current": round(current_score, 3),
            "previous": round(previous_score, 3),
            "drop": round(delta, 3),
        }
    return {"status": "OK", "current": round(current_score, 3)}
```

What to Monitor

Metric	How to Measure	Alert Threshold
Pass rate	Golden dataset test runs	Drop > 5% from baseline
Average quality score	LLM-as-judge average	Drop below 0.85
Bias delta	Sentiment gap across demographic groups	Gap > 10%
Category regression	Per-category pass rates	Any category drops > 15%
Response latency	P95 response time	Exceeds 5 seconds

The Monitoring Schedule

Run evaluations at three cadences:

- **On every prompt change:** Full golden dataset test before deploying prompt updates.
- **Weekly:** Automated regression test to catch drift from provider model updates.
- **Monthly:** Full bias audit with expanded demographic test suite and manual review of flagged cases.

⚠ Warning

A passing evaluation today does not guarantee a passing evaluation next month. Model providers update their models regularly. OpenAI, Anthropic, and Google all push model updates that can subtly change output behavior. Your code did not change, but the model behind the API did. Regular re-evaluation is not optional.

Try It Yourself

Exercise 1: Run a Bias Swap Test

Use the `test_for_bias` function with a prompt template of your choice. Test with at least three different names from different cultural backgrounds. Compare the outputs and note any differences in tone, specificity, or recommendation quality.

Exercise 2: Build a Mini Evaluation Harness

Create a golden dataset with 5 test cases relevant to your application. Implement keyword-based scoring and run all 5 cases. Print a summary showing pass rate and the worst-performing test case.

Exercise 3: Add Drift Detection

Run your evaluation harness twice with different prompts (one concise, one verbose). Use the `check_for_drift` function to compare results and verify it flags the regression if one prompt performs worse.

Tip

Starter Code: The companion repository contains full exercises, starter code, and solutions for building an evaluation harness to detect regressions and measure bias.

- [building-with-llms-companion/exercises/ch14/eval_harness.py](https://github.com/Kangsk/Kangsk.github.io/blob/main/building-with-llms-companion/exercises/ch14/eval_harness.py)

Chapter Summary

Key Takeaways

- Bias enters through training data, labeling, prompt design, and evaluation gaps. The swap test (same prompt, different names or backgrounds) is the fastest way to surface it. Mitigate with fairness instructions, diverse few-shot examples, post-processing filters, and regular audits run after every prompt change.
- In medical, legal, financial, and hiring domains, LLMs should assist human decisions rather than replace them. Match the level of human oversight to risk: automate low-stakes tasks, augment medium-risk ones, and require human approval for high-stakes decisions.
- Build an evaluation harness with 50 or more golden test cases, combine keyword scoring with LLM-as-judge, and run drift detection weekly. Model providers push updates silently, and behavior can degrade without any change to your code.

Common Pitfalls

- Assuming "be fair and unbiased" in a system prompt eliminates bias (it helps but is not sufficient)
- Using the same model as both generator and judge (self-evaluation is unreliable)
- Running evaluation once at launch and never again (drift happens silently)

Knowledge Check

1. Multiple Choice: What is training data bias?

☐ Models intentionally discriminating against groups

☐ Models learning prejudices from patterns in their training data

☐ A bug in the model's code

☐ Biased API pricing

2. **True or False:** Adding "Be fair and unbiased" to the system message is sufficient to eliminate bias in LLM outputs.

☐ True ☐ False

3. **Fill in the Blank:** A ____-in-the-loop design pattern requires a human to approve AI decisions before they are enacted.

4. **Multiple Choice:** What is drift detection?

☐ Debugging code after deployment

☐ Monitoring model accuracy over time to catch gradual degradation

☐ Detecting prompt injection attacks

☐ Tracking API cost changes

5. **Scenario:** Your resume screener rates candidates differently based on name origin. "James Smith" receives "strong candidate" while "Lakisha Washington" (with identical qualifications) receives "needs further review." Name three mitigation steps you would implement.

Answers

1. **(b) Models learning prejudices from patterns in their training data.**

LLMs learn statistical patterns from internet text, which reflects societal biases. The model is not intentionally biased; it reproduces patterns it learned during training.

2. **False.** Fairness instructions in the system message reduce bias but do not eliminate it. Effective mitigation requires multiple layers: debiased prompts, diverse few-shot examples, post-processing filters, and regular bias audits with demographic swap tests.

3. **Human-in-the-loop.** In this pattern, the AI generates recommendations or drafts, but a human makes the final decision. The level of human involvement scales with the risk level of the decision.

4. **(b) Monitoring model accuracy over time to catch gradual degradation.** Drift occurs when model behavior changes due to provider updates, shifting user patterns, or evolving data. Regular re-

evaluation against a golden dataset is the primary detection mechanism.

5. Three mitigation steps: **(1)** Remove or blind names in prompts sent to the LLM so the model evaluates qualifications only. **(2)** Add explicit fairness instructions to the system prompt: "Evaluate based solely on skills, experience, and qualifications." **(3)** Build a bias test suite with diverse names and run it after every prompt change. Additional steps include adding human review for all hiring recommendations and auditing past decisions for systematic bias. Cross-reference [Chapter 12](#) for security guardrails that complement ethical guardrails.