

Appendix D: Further Reading & Resources

Papers

- Attention Is All You Need (Vaswani et al., 2017) - The foundational Transformer paper
- Language Models are Few-Shot Learners (Brown et al., 2020) - GPT-3 and in-context learning
- Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks (Lewis et al., 2020)
- Chain-of-Thought Prompting Elicits Reasoning in Large Language Models (Wei et al., 2022)

Documentation

- OpenAI API Documentation: <https://platform.openai.com/docs>
- Anthropic Claude Documentation: <https://docs.anthropic.com>
- Google Gemini Documentation: <https://ai.google.dev/docs>
- LiteLLM Documentation: <https://docs.litellm.ai>

Communities

- Hugging Face: <https://huggingface.co>
- OpenAI Developer Forum: <https://community.openai.com>
- r/MachineLearning: <https://reddit.com/r/MachineLearning>

Books & Courses